

Dental Vision Benchmark

Which model actually reads a dental image?

1 · IMAGES

Open-licensed (Wikimedia, open access)
+ one transformed control item (flip / crop)

captions removed

2 · VISION CONTESTANTS

image only · no caption · no internet

GPT-5.5

Claude

Gemini

Qwen-VL

Llama / GLM

each writes a free-text description

GROUND TRUTH

Clinician writes, per image:

must_identify

must_avoid

+ reference caption

You sign off.

3 · TWO JUDGES

independent vendors grade against the truth:
covered · violations · modality · agreement

each sees image + the real answer key

4 · RESULTS

- Leaderboard — the best dental-image reader, with confidence intervals
- Original vs transformed control — a recognition sanity check
- Task split — diagrams readable, unlabeled clinical photos hard

□ teal = the two safeguards: no leakage going in, the real answer key at judging